

CSCI-3160: COMPUTER SYSTEMS

STATICALLY-SCHEDULED PROCESSOR PIPELINE SIMULATION

GOAL

Your team will create a simulator that accepts, as input, a stream of 32-bit RISC-V instructions, and outputs a description of when each of those instructions passes through different stages of an in-order processor's pipeline. Your team, at project completion, will present your simulation to the class.

SPECIFICATIONS

Keep in mind: these specifications are the “bare minimum” I expect for this assignment, i.e., a complete and correct implementation of these deliverables (the revised instruction set—if needed, the in-order pipeline simulation, and the presentation) are what I would expect a ‘C’-level effort to look like. Students with more ambitious designs and/or extensions of this specification may be able to use that additional effort and work as evidence for a higher grade at the end of the course.

INSTRUCTION SET

We will use the canonical RISC-V instruction set architecture for the simulation:

- Data transfers: `lw`, `sw`
- Arithmetic: `add`, `addi`, `sub`
- Control: `<nop` is a pseudo-instruction for `addi x0, x0, 0>`, `bne`, `beq`
- Register file: `x0–x31`

PIPELINE DESIGN

At a minimum: your project will simulate the RISC-V instruction subset specified above in a five-stage pipeline:

1. Fetch instruction
2. Decode instruction / read register file
3. Execute instruction / calculate address
4. Access memory
5. Write result to register file

ASSUMPTIONS AND LATENCIES

Not all instructions take the same amount of time to execute. However, for the purposes of this (minimum) specification, I'm making some assumptions about latencies of these instructions. For your own instruction set, you may have similar assumptions.

- At most, one instruction can be issued (fetched) per cycle.
- At most, one data memory access (reading/writing memory) can occur per cycle.
 - Memory accesses take three cycles (“magic” L1 cache with 100% hit rate)
- Loads and stores require one cycle to calculate the effective address
- The register file may **not** accomplish a read **and** a write of the **same** register in the same cycle
 - For simplicity's sake (not particularly realistic)
- **Only one instruction may be executing (in the ALU) at a time:**
 - Integer arithmetic (including address calculations) complete in one cycle.

Since this base specification only simulates the pipeline (and not the actual instruction effects), I assume that addresses given in different registers (along with offsets) will not ever collide, e.g., adding an offset to register `x3` will **not** lead to the same base address held in `x2`.

Finally: if I encounter any structural (e.g., memory access taking more than one cycle), data hazards (due to a RAW, WAR, or WAW dependency), or branching hazards (due to not knowing if the branch will be taken or fallen through) in the processor pipeline, I stall (insert a bubble into the pipeline).

OTHER DETAILS

Anything not specified above is left to your discretion—however, be prepared to justify your decisions. See me if you want feedback on your ideas before implementing them.

DELIVERABLES

Your team will prepare a short presentation and demonstration of your work at final project delivery during class (15 minutes). Your presentation should also include a short project retrospective about who worked on what, how the project work went, and any improvements to process you made / would make on a future project.

EXTENSIONS

This project has a lot of room for some interesting extensions—primarily, forwarding via pipeline registers (instead of stalling) on data hazards, static branch prediction paired with flushing (instead of stalling) for control hazards, implementing control logic considerations (pairs well with forwarding), and building a GUI datapath for your pipeline to illustrate what's happening at each cycle.

These are just a handful of ideas you and your team could implement. I am open to other project extensions (like, say, static multi-issue) or implementing a dynamically-scheduled pipeline (which opens the door to the really interesting instruction-level parallelism features and interesting floating-point instruction delays).

EXAMPLE INPUTS AND OUTPUTS

CONFIG

Notes on this file: I use this configuration file format across several of my simulations—this is **not** a hard-and-fast format but merely an example. You may hard-code your configurations.

For my buffers section: if I only have a total of one arithmetic buffer, I assume the “magic” ALU/FPU combo described above in the assumptions.

```
buffers

fetch: 1
memory: 1
ints: 1
fp adds: 0
fp muls: 0
```

```
latencies

memory: 3
ints: 1
fp_add: 2
fp_mul: 5
fp_div: 10
```

TRACE.DAT (INPUT INSTRUCTIONS)

Notes on this file: I assume no data races in main memory on address resolution (i.e., no collisions). Your simulation may accept assembly language instructions or machine language instructions (or both—just in separate files).

```
flw f6,32(x2)
flw f2,48(x3)
fmul.s f0,f2,f4
fsub.s f8,f6,f2
fdiv.s f10,f0,f6
fadd.s f6,f8,f2
```

OUTPUT

Configuration

buffers:

```
fetch: 1
memory: 1
ints: 1
fp adds: 0
fp muls: 0
```

NOTE combining ALU with FPU

Latencies:

```
memory: 3
ints: 1
fp_add: 2
fp_mul: 5
fp_div: 10
```

Static Pipeline Simulation

Instruction	Inst. Fetch	Decode Read Reg	Execute Calc Adr	Access Memory	Write Register
flw f6, 32(x2)	1	2	3	4 - 6	7
flw f2, 48(x3)	2	3	4 - 6	7 - 9	10
fmul.s f0, f2, f4	3	4 - 11	12 - 16		17
fsub.s f8, f6, f2	4 - 11	12 - 16	17 - 18		19
fdiv.s f10, f0, f6	12 - 16	17 - 18	19 - 28		29
fadd.s f6, f8, f2	17 - 18	19 - 28	29 - 30		31

Hazards

structural: 7
data: 2
control: 0

Dependencies

read-after-write: 2
write-after-read: 0
write-after-write: 0

Cycles Stalled

instruction fetch: 12
inst. decode / read register: 21
execute / calculate address: 2
read / write memory: 0
write register: 0

Notes on this output: the hazards statistics is counting the number of times my simulation inserted bubbles into the pipeline—some of these overlap (like the decode and read register step for the floating-point multiply instruction—I had a structural hazard due to the stall in the execution stage **and** a data hazard due to the read-after-write dependency on f2).

The dependencies are only those that result in an actual hazard—there's 7 occurrences of a register being read after a write, but only 2 of those occurrences would overlap temporally to cause a stale read. In addition, there's 1 occurrence of a write (f6) after a read, but by the time the floating-point addition would write its result to f6, the reads of f6 are complete (thus, no WAR dependency).